

Lab 8: Throughput Engineering — Bulk Import, Batching & the Schema Pattern Playbook (75 min)

This is the **performance capstone of Day 2**. Everything you measure here is a number you can quote in a design review. Pairs with the [Schema Patterns Playbook](#).

Learning Objectives

By the end of this lab you will be able to:

- Measure and rank the four ways of getting rows into CockroachDB: single-row `INSERT`, multi-row `INSERT`, `COPY`, and `IMPORT INTO`
- Find the **batch-size knee** for your workload instead of guessing at 100 or 1000
- Pre-split a table with `ALTER TABLE ... SPLIT AT` and prove the import got faster
- Run the same write workload against four PK designs and produce a rows/sec table
- Replace a single-row counter with a **Sharded Counter** and measure the retry rate collapse
- Run an online schema change under live write load and quantify the impact
- Use `cockroach workload` to find the concurrency knee where throughput plateaus but p99 explodes
- Size a connection pool from measured numbers rather than folklore

Prerequisites

- **Docker Desktop** (or Docker Engine) running — there is no `cockroach` binary to install
- Labs 3 (schema design) and 6 (`EXPLAIN ANALYZE`) — you'll reuse both
- `psql` on `PATH` for the `COPY` part (optional; a fallback is provided)

Setup

Start a 3-node demo cluster. Three nodes is the minimum that makes distribution visible.

```
scripts/crdb up          # start the 3-node cluster (skip if it is already running)
scripts/crdb sql         # open a SQL shell
```

Everything runs in Docker — see [Lab 1](#) for the cluster layout. Use `scripts/crdb.sh` on macOS, Linux, and **WSL** — WSL is Linux, so the shell script is the right one there. Only a native Windows shell (`cmd` or PowerShell) needs `scripts\crdb.bat`.

```
CREATE DATABASE throughput;
USE throughput;
SET sql_safe_updates = off;
\timing on
```

Several parts of this lab drive the cluster from a **second terminal**. Copy the connection URL from the demo banner (the line starting `postgresql://demo@127.0.0.1:...`) and export it there:

```
export CRDB_URL='postgresql://root@localhost:26257/?sslmode=disable'
scripts/crdb sql -e "SELECT 1;" # confirm before continuing
```

Inside the demo shell, `\demo ls` reprints the connection parameters for every node if you lose the banner.

Keep a results table open. Every part of this lab produces a number. Record it — the final deliverable is the filled-in table at the end, not the individual commands.

Tasks

Part A: Four Ways to Load 100,000 Rows (15 min)

The target schema is deliberately boring so the measurement is about the *method*, not the design:

```
CREATE TABLE load_test (
  id          UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  tenant      INT NOT NULL,
  amount      DECIMAL(12,2) NOT NULL,
  note        STRING NOT NULL,
  created     TIMESTAMPTZ NOT NULL DEFAULT now()
);
```

A1. Single-row inserts (*the baseline you never want in production*)

Because a per-row round trip is the point, run this from the shell, not from SQL:

```
# In a second terminal. --url comes from the demo banner (postgresql://...).
time (for i in $(seq 1 2000); do
  echo "INSERT INTO throughput.load_test (tenant, amount, note) VALUES (1, 9.99, 'row');"
done | scripts/crdb sql)
```

2000 rows — not 100,000. Extrapolate:

```
-- rows/sec for the single-row method
SELECT 2000 / <seconds_from_time> AS rows_per_sec;
```

Why so slow? Each statement is its own implicit transaction: a full Raft round trip per row, plus a network round trip per row. You are measuring latency, not throughput.

A2. Multi-row INSERT (batched)

```
TRUNCATE load_test;

-- 100,000 rows as 100 statements of 1,000 rows, generated server-side
INSERT INTO load_test (tenant, amount, note)
SELECT (g % 50), (g % 1000)::DECIMAL / 100, 'batched row ' || g
FROM generate_series(1, 100000) g;
```

Record the elapsed time. This is one statement, one transaction, and CockroachDB pipelines the writes across ranges.

A3. COPY

```
# Generate a CSV once
python3 - <<'PY'
import csv, random, uuid
with open('/tmp/load_test.csv', 'w', newline='') as f:
    w = csv.writer(f)
    for i in range(100_000):
        w.writerow([str(uuid.uuid4()), i % 50, round(random.random()*1000, 2),
                    f'copy row {i}', '2024-01-01 00:00:00+00'])
PY

time psql "$CRDB_URL" -c "TRUNCATE throughput.load_test" \
-c "\copy throughput.load_test (id, tenant, amount, note, created) FROM '/tmp/
load_test.csv' CSV"
```

No psql? Use `scripts/crdb sql -e "\copy ..."` — the built-in shell supports `\copy` with the same syntax.

A4. IMPORT INTO

`IMPORT INTO` bypasses the SQL layer entirely and writes SSTables directly into the storage layer. It takes the table **offline** for the duration.

```
TRUNCATE load_test;
```

```
# The CSV is on your machine; the node is in a container and cannot see it.
# Copy it in first, then upload it to the cluster's userfile store.
scripts/crdb cp /tmp/load_test.csv crdb1:/tmp/load_test.csv
scripts/crdb run userfile upload /tmp/load_test.csv /lab8/load_test.csv --insecure
```

This two-step is the containerised version of "put the file where the database can read it". In production the same problem is solved by cloud storage — `IMPORT INTO ... CSV DATA ('s3://...')` — which is why real bulk loads point at S3/GCS rather than a local path.

```
IMPORT INTO load_test (id, tenant, amount, note, created)
CSV DATA ('userfile:///lab8/load_test.csv');
```

A5. Record the results

Method	Rows	Elapsed	Rows/sec	Table online during load?
Single-row <code>INSERT</code>	2,000 (extrapolated)			☐
Multi-row <code>INSERT</code>	100,000			☐
<code>COPY</code>	100,000			☐
<code>IMPORT INTO</code>	100,000			☐ offline

The decision rule. `IMPORT INTO` for one-time loads and migrations where downtime on that table is acceptable. `COPY` for large loads that must stay online. Multi-row `INSERT` for application-driven ingest. Single-row `INSERT` only when the row arrives alone.

Part B: Find the Batch-Size Knee (10 min)

"Batch your inserts" is useless advice without a number. Find yours.

```
TRUNCATE load_test;

-- Run each of these and record the elapsed time.
-- Total rows is constant (50,000); only the rows-per-statement changes.

-- batch = 1000 (50 statements)
INSERT INTO load_test (tenant, amount, note)
SELECT g % 50, 1.00, 'b' FROM generate_series(1, 1000) g; -- repeat 50x

-- Easier: drive the sweep from one script
```

```

for BATCH in 1 10 100 1000 5000 10000; do
  STMTS=$((50000 / BATCH))
  START=$(date +%s.%N)
  for i in $(seq 1 $STMTS); do
    echo "INSERT INTO throughput.load_test (tenant, amount, note) SELECT g % 50, 1.00, 'b'
        FROM generate_series(1, $BATCH) g;"
  done | scripts/crdb sql >/dev/null 2>&1
  END=$(date +%s.%N)
  echo "batch=$BATCH elapsed=$(echo "$END - $START" | bc)s rows/sec=$(echo "50000 / ($END
    - $START)" | bc)"
  scripts/crdb sql -e "TRUNCATE throughput.load_test" >/dev/null
done

```

Batch size	Elapsed	Rows/sec	Notes
1			latency-bound
10			
100			
1,000			usually near the knee
5,000			
10,000			watch for txn size limits

What you should see: steep gains from 1 → 100, a knee somewhere around 500–2,000, then flat or *worse* past that. Very large batches hold more intents open, raise contention and retry cost, and risk hitting `kv.transaction.max_intents_bytes`. When a batch fails, you retry the *whole* batch — the knee is about failure blast radius as much as it is about speed.

Part C: Pre-Split Before a Bulk Load (*Playbook #10*) (10 min)

A new table is one range. A 100M-row import into one range means every write goes to one leaseholder until the splitter catches up.

1. Un-split baseline:

```

CREATE TABLE seq_import (id INT PRIMARY KEY, payload STRING);

SELECT count(*) AS ranges_before FROM [SHOW RANGES FROM TABLE seq_import];

INSERT INTO seq_import
SELECT g, repeat('x', 200) FROM generate_series(1, 200000) g;

SELECT count(*) AS ranges_after FROM [SHOW RANGES FROM TABLE seq_import];

```

Record the elapsed time and the range count.

2. Pre-split version:

```
CREATE TABLE seq_import_split (id INT PRIMARY KEY, payload STRING);

ALTER TABLE seq_import_split SPLIT AT
  SELECT g * 10000 FROM generate_series(1, 19) g;

SELECT count(*) AS ranges_before FROM [SHOW RANGES FROM TABLE seq_import_split];

INSERT INTO seq_import_split
SELECT g, repeat('x', 200) FROM generate_series(1, 200000) g;
```

3. Compare leaseholder spread:

```
SELECT lease_holder, count(*) AS ranges
FROM [SHOW RANGES FROM TABLE seq_import_split WITH DETAILS]
GROUP BY lease_holder ORDER BY lease_holder;
```

4. Release the manual splits when the load is done — otherwise the splits are pinned forever and the cluster can't merge them back:

```
ALTER TABLE seq_import_split UNSPLIT ALL;
```

Sizing the split points. One split per ~512 MiB of expected data, or one per node × a small multiple (3–10), whichever is larger. You are seeding the allocator, not micro-managing it.

Part D: Four PK Designs Under Concurrent Write Load (15 min)

Lab 3 showed you the *shape* of a hotspot. Now put a number on it.

1. Create the four tables:

```
CREATE TABLE pk_serial (
  id SERIAL PRIMARY KEY, tenant INT, payload STRING, created TIMESTAMPTZ DEFAULT now());

CREATE TABLE pk_uuid (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(), tenant INT, payload STRING,
  created TIMESTAMPTZ DEFAULT now());

CREATE TABLE pk_composite (
  tenant INT, id UUID DEFAULT gen_random_uuid(), payload STRING,
  created TIMESTAMPTZ DEFAULT now(), PRIMARY KEY (tenant, id));

CREATE TABLE pk_hash (
  tenant INT, created TIMESTAMPTZ DEFAULT now(), id UUID DEFAULT gen_random_uuid(),
  payload STRING,
  PRIMARY KEY (tenant, created, id) USING HASH WITH (bucket_count = 16));
```

2. Drive 8 concurrent writers at each table and time the whole run:

```

run_writers() {
  TABLE=$1; ROWS_PER_WRITER=2000; WRITERS=8
  START=$(date +%s.%N)
  for w in $(seq 1 $WRITERS); do
    ( for i in $(seq 1 $((ROWS_PER_WRITER / 100))); do
      echo "INSERT INTO throughput.$TABLE (tenant, payload) SELECT g % 50, 'p' FROM
      generate_series(1,100) g;"
    done | scripts/crdb sql >/dev/null 2>&1 ) &
  done
  wait
  END=$(date +%s.%N)
  TOTAL=$((ROWS_PER_WRITER * WRITERS))
  echo "$TABLE: $TOTAL rows in $(echo "$END - $START" | bc)s => $(echo "$TOTAL / ($END -
  $START)" | bc) rows/sec"
}

for t in pk_serial pk_uuid pk_composite pk_hash; do run_writers $t; done

```

3. Confirm the distribution difference:

```

SELECT 'pk_serial' AS t, count(*) AS ranges FROM [SHOW RANGES FROM TABLE pk_serial]
UNION ALL SELECT 'pk_uuid', count(*) FROM [SHOW RANGES FROM TABLE pk_uuid]
UNION ALL SELECT 'pk_composite', count(*) FROM [SHOW RANGES FROM TABLE pk_composite]
UNION ALL SELECT 'pk_hash', count(*) FROM [SHOW RANGES FROM TABLE pk_hash];

```

4. Check the hot-range signature while a run is in flight:

```

-- SQL gives you distribution, not traffic:
SELECT table_name, count(*) AS ranges, count(DISTINCT lease_holder) AS leaseholders
FROM [SHOW CLUSTER RANGES WITH TABLES, DETAILS]
WHERE database_name = 'throughput'
GROUP BY table_name ORDER BY ranges DESC;

```

For **traffic** per range, use **DB Console** → **Advanced Debug** → **Hot Ranges** (or the `/_status/hotranges` endpoint) — per-range QPS is not exposed through SQL.

5. Now measure the cost of the design. Hash-sharding is not free — it makes ordered scans fan out across every bucket:

```

EXPLAIN ANALYZE
SELECT * FROM pk_hash
WHERE tenant = 7 AND created > now() - INTERVAL '1 hour'
ORDER BY created DESC LIMIT 100;

EXPLAIN ANALYZE
SELECT * FROM pk_composite WHERE tenant = 7 ORDER BY id LIMIT 100;

```

Compare the **rows read** and the number of spans scanned in each plan.

PK design	Write rows/sec	Ranges	Ordered-scan cost	Use when
SERIAL			cheapest	never on a hot path
UUID			no useful order	writes dominate, point reads
(tenant, id)			ordered per tenant	multi-tenant, tenant-scoped queries
hash-sharded			fans out over buckets	time-ordered writes at high rate

Part E: Sharded Counter vs Single-Row Counter (*Playbook #5*) (10 min)

1. The anti-pattern:

```
CREATE TABLE counter_single (name STRING PRIMARY KEY, n INT NOT NULL DEFAULT 0);
INSERT INTO counter_single VALUES ('page_views', 0);
```

2. The pattern:

```
CREATE TABLE counter_shards (
  name STRING,
  shard INT2,
  n INT NOT NULL DEFAULT 0,
  PRIMARY KEY (name, shard)
);
INSERT INTO counter_shards (name, shard, n)
SELECT 'page_views', g, 0 FROM generate_series(0, 15) g;
```

3. Record the retry baseline before you start:

```
SELECT sum((statistics->'statistics'->'maxRetries')::INT) AS max_retries
FROM crdb_internal.statement_statistics;
```

4. First, the trap — pick the shard the obvious way and count the result:

```
UPDATE counter_shards SET n = 0 WHERE name = 'page_views';
```

```
# 200 sequential increments, no concurrency at all
for i in $(seq 1 200); do
  echo "UPDATE throughput.counter_shards SET n = n + 1
    WHERE name = 'page_views' AND shard = (random()*16)::INT2;"
done | scripts/crdb sql >/dev/null 2>&1
```

```
SELECT sum(n) AS should_be_200 FROM counter_shards WHERE name = 'page_views';
```

It will not be 200. `random()` is a **volatile** function, and a volatile function in a `WHERE` predicate is evaluated **once per row scanned** — not once per statement. The statement therefore matches a random *number* of the 16 shard rows: often one, sometimes zero (the increment vanishes), sometimes three (it counts triple). See it directly:

```
SELECT count(*) FROM counter_shards WHERE shard = (random()*16)::INT2; -- run 8 times
```

You'll get a different count nearly every time. A scalar subquery (`shard = (SELECT (random()*16)::INT2)`) is no better.

The fix: choose the shard where you have a stable value — in the application — and bind it as a parameter. This is the single most common way a sharded counter is written wrong, and because the loss is silent it survives code review.

5. Now do it correctly, and hammer both designs with 16 concurrent bumpers:

```
bump_single() {
    START=$(date +%s.%N)
    for w in $(seq 1 16); do
        ( for i in $(seq 1 100); do
            echo "UPDATE throughput.counter_single SET n = n + 1 WHERE name = 'page_views';"
            done | scripts/crdb sql >/dev/null 2>&1 ) &
        done
    wait
    echo "single: 1600 increments in $(echo "$(date +%s.%N) - $START" | bc)s"
}

bump_sharded() {
    START=$(date +%s.%N)
    for w in $(seq 1 16); do
        ( for i in $(seq 1 100); do
            # the SHELL picks the shard – a literal by the time the server sees it
            echo "UPDATE throughput.counter_shards SET n = n + 1
                WHERE name = 'page_views' AND shard = $((RANDOM % 16));"
            done | scripts/crdb sql >/dev/null 2>&1 ) &
        done
    wait
    echo "sharded: 1600 increments in $(echo "$(date +%s.%N) - $START" | bc)s"
}

scripts/crdb sql -e \
    "UPDATE throughput.counter_single SET n = 0; UPDATE throughput.counter_shards SET n =
    0;"
bump_single
bump_sharded
```

Verify both counted every increment before you compare the times — a fast wrong answer is not a win:

```
SELECT (SELECT n FROM counter_single WHERE name = 'page_views') AS
       single_total,
       (SELECT sum(n) FROM counter_shards WHERE name = 'page_views') AS
       sharded_total;
-- both must equal 1600
```

6. Count the retries each design generated:

```
SELECT
  substring(metadata->'query', 1, 60) AS stmt,
  (statistics->'statistics'->'cnt')::INT AS executions,
  (statistics->'statistics'->'maxRetries')::INT AS max_retries,
  round((statistics->'statistics'->'svcLat'->'mean')::FLOAT * 1000, 2) AS mean_ms
FROM crdb_internal.statement_statistics
WHERE metadata->'query' LIKE '%counter%'
ORDER BY executions DESC;
```

7. Read the counter:

```
SELECT sum(n) FROM counter_shards WHERE name = 'page_views';
```

The read is now a 16-row scan instead of a 1-row point lookup. That's the trade: reads get marginally more expensive, writes stop serializing.

Choosing the shard count. Start at your peak concurrent writer count, round to a power of two, cap around 64. Too few shards leaves contention; too many makes the read scan wide and spreads the rows over more ranges than the workload needs.

Part F: Online Schema Change Under Load (8 min)

1. Start a background write stream:

```
( while true; do
  echo "INSERT INTO throughput.pk_uuid (tenant, payload) SELECT g % 50, 'live' FROM
    generate_series(1,200) g;"
done | scripts/crdb sql >/dev/null 2>&1 ) &
WRITER_PID=$!
```

2. Add an index while writes are running:

```
CREATE INDEX ON pk_uuid (tenant, created DESC);
```

3. Watch it as a job, not as a lock:

```
SELECT job_id, job_type, description, status, fraction_completed
FROM [SHOW JOBS] WHERE job_type IN ('SCHEMA CHANGE', 'NEW SCHEMA CHANGE')
ORDER BY created DESC LIMIT 5;
```

Two schema-changer job types. Modern CockroachDB runs most DDL through the *declarative* schema changer, which records jobs as `NEW SCHEMA CHANGE`. The legacy changer still handles some operations (e.g. `TRUNCATE`) and uses `SCHEMA CHANGE`. Filtering on only one of them is why people conclude "my `CREATE INDEX` didn't create a job" — match both.

4. Measure the write-path cost of the new index:

```
SELECT
  (statistics->'statistics'->>'cnt')::INT AS executions,
  round((statistics->'statistics'->'svcLat'->>'mean')::FLOAT * 1000, 2) AS mean_ms
FROM crdb_internal.statement_statistics
WHERE metadata->'query' LIKE 'INSERT INTO pk_uuid%';
```

Every secondary index multiplies write amplification. Two indexes on a table means each `INSERT` writes three key-value entries, and each one is a Raft proposal.

5. Stop the writer:

```
kill $WRITER_PID 2>/dev/null; pkill -f "generate_series(1,200)" 2>/dev/null
```

The rule that matters in a deploy: `CREATE INDEX`, `ADD COLUMN` with a default, and `DROP COLUMN` run as **background jobs** — non-blocking, resumable, and cancellable while traffic continues. `ALTER COLUMN TYPE` and adding a `NOT NULL` constraint to existing data require a validating backfill and are the ones to schedule carefully.

Part G: The Concurrency Knee — `cockroach workload` (7 min)

Throughput plateaus long before latency does. Find the point where adding connections only adds queueing.

```
scripts/crdb run workload init kv --drop 'postgresql://root@crdb1:26257?sslmode=disable'

for C in 1 4 16 64 256; do
  echo "=== concurrency $C ==="
  scripts/crdb run workload run kv \
    --duration=30s --concurrency=$C \
    --read-percent=50 --max-rate=0 \
    "$CRDB_URL" 2>&1 | tail -4
done
```

Concurrency	Throughput (ops/s)	p50 (ms)	p99 (ms)
1			
4			
16			
64			
256			

Plot throughput and p99 against concurrency. The knee is the last concurrency level where throughput still climbs meaningfully. Past it you are buying latency with no throughput.

Pool sizing from these numbers. Little's Law: $\text{concurrency} = \text{throughput} \times \text{latency}$. If the knee is 64 concurrent statements at 8 ms mean latency, a single app instance sustaining 2,000 QPS needs $2000 \times 0.008 \approx 16$ connections — not 200. Size the pool to the knee divided by the number of app instances, then add a small headroom margin. An oversized pool does not add throughput; it moves the queue from your app into the database, where it is harder to see and more expensive to drain.

Results Summary — Fill This In

Measurement	Your number	Rule of thumb
Single-row insert rows/sec		latency-bound, $\sim 1/\text{RTT}$
Multi-row insert rows/sec		10–100× single-row
<code>IMPORT INTO</code> rows/sec		fastest, table offline
Batch-size knee		usually 500–2,000
Pre-split speedup		biggest on sequential keys
Best PK design (writes)		hash-sharded or UUID
Worst PK design (writes)		<code>SERIAL</code>
Sharded counter speedup		grows with concurrency
Concurrency knee		pool size starts here

Cleanup

```
DROP DATABASE throughput CASCADE;
DROP DATABASE IF EXISTS kv CASCADE;
```

```
rm -f /tmp/load_test.csv
```

Then `\q` to exit the demo cluster.

The cluster keeps running between labs — that is the point of it being persistent. To wipe everything and start fresh at any time:

```
scripts/crdb reset
```

Lab 8 Deliverables

□ **Ingest ranking**: four load methods measured, with a stated decision rule for each □ **Batch knee**: an actual number for this schema, not a guess □ **Pre-split**: range count and load time before/after □ **PK bake-off**: rows/sec for four designs plus the ordered-scan cost each one pays □ **Sharded counter**: retry count and throughput vs the single-row version, *and* the volatile-predicate trap observed and fixed □ **Online schema change**: index added under live write load, job observed, write cost measured □ **Concurrency knee**: throughput/p99 curve and a pool size derived from it

Challenge Exercises

1. **Beat IMPORT INTO**. Pre-split `load_test` on the UUID keyspace before the import (`SPLIT AT VALUES ('11111111-...'), ('22222222-...'), ...`). Does the import get faster? Why is the effect smaller than it was for the sequential-key table?
2. **Find the write-amplification curve**. Add indexes to `pk_uuid` one at a time (0, 1, 2, 4) and re-measure insert throughput at each step. Plot rows/sec against index count. What is the marginal cost of index number four?
3. **Break the batch**. Set `SET CLUSTER SETTING kv.transaction.max_intents_bytes = '1MiB'`, then find the batch size that starts failing. What error do you get, and what should the application do with it?
4. **Re-run Part D on more nodes**. Add `crdb5 ... crdb9` to `docker/labs.yml` (copy the `crdb4` block, bump the ports) and re-run. Which PK design's advantage grows with node count, and which one's shrinks?

Reference

Command	Purpose
<code>IMPORT INTO t (cols) CSV DATA ('...')</code>	Fastest bulk load; table goes offline
<code>\copy t FROM 'file' CSV</code>	Fast bulk load; table stays online
<code>ALTER TABLE t SPLIT AT SELECT ...</code>	Seed ranges before a bulk load
<code>ALTER TABLE t UNSPLIT ALL</code>	Release manual splits so ranges can merge
<code>scripts/crdb cp <file> crdb1:/tmp/</code>	Copy a host file into a node container
<code>scripts/crdb run userfile upload</code>	Stage a file in the cluster for <code>IMPORT</code>
<code>scripts/crdb run workload run kv --concurrency=N</code>	Throughput/latency curve
<code>crdb_internal.statement_statistics</code>	Executions, retries, mean latency per statement
Volatile functions (<code>random()</code> , <code>now()</code>) in a predicate	Re-evaluated per row — never use them to select a row
<code>SHOW RANGES FROM TABLE t WITH DETAILS</code>	Range count, size, leaseholder placement
<code>SHOW JOBS</code>	Watch online schema changes progress