

Lab 11: BACKUP / RESTORE / Schedules & Cross-Cluster DR Drill (65 min)

Learning Objectives

By the end of this lab you will be able to:

- Take full, incremental, and scheduled backups and read backup metadata without restoring
- Restore a database, a single table, and a point in time — into a new name, side by side
- Run a **cross-cluster DR drill**: back up from cluster A, restore into cluster B, verify
- Compute the RTO and RPO your backup cadence actually guarantees, and measure the real RTO
- Use `SHOW BACKUP` and revision history to answer "what did this table look like at 14:02?"
- Write a DR runbook with the specific commands your on-call will run at 3am

Prerequisites

- **Docker Desktop** (or Docker Engine) running — there is no `cockroach` binary to install
- ~4 GB free RAM (two small clusters run side by side in Part D)

Setup — Cluster A (primary)

```
scripts/crdb up
export A='postgres://root@localhost:26257?sslmode=disable'
```

The cluster runs in Docker (see [Lab 1](#)). From your machine it is `localhost:26257`; from inside another container it is `crdb1:26257`. `scripts/crdb run ...` executes inside node 1.

Backup storage in this lab. We use a local `node1:26257` path so the lab runs offline. In production the destination is cloud storage: `s3://bucket/path?AUTH=implicit`, `gs://bucket/path?AUTH=implicit`, or `azure-blob://...`. The SQL is identical apart from the URI.

⚠ Licensing — read this before Part A

On a cluster with **no enterprise licence**, the backup features split like this (verified on v23.2). A licence is free for training — see option 3 below — but the lab is written so that every lesson lands without one:

Feature	Without a license
Full BACKUP / RESTORE	☐ works
RESTORE ... WITH new_db_name / into_db	☐ works
Full-cluster BACKUP / RESTORE (the DR drill)	☐ works
CREATE SCHEDULE	☐ works — full backups only
WITH revision_history	☐ enterprise
Incremental (INTO LATEST IN ...)	☐ enterprise
Point-in-time RESTORE ... AS OF SYSTEM TIME	☐ enterprise (needs revision history)

The error is explicit, e.g.: `use of BACKUP with revision_history requires an enterprise license`.

Three ways to run this lab:

1. **Free path (default).** Do every step; where a step is marked ☐Enterprise, read it and run the free alternative given beside it. You still complete the DR drill end to end.
2. **A containerised demo cluster.**
`docker run --rm -it cockroachdb/cockroach:v23.2.5 demo --nodes 3 --no-example-database --empty` ships with a temporary licence, so every step works. Use it for Parts A–C; Part D needs two clusters, so start a second demo on other ports.
3. **A licence — free for training, and the class has one.** It ships with the repo in `.license.env`; `scripts/crdb up` prints `enterprise licence applied` when it starts, and nothing further is needed. If your cluster was started before the file existed, run `scripts/crdb reset` — or apply the key from that file in place:

```
-- paste the value of COCKROACH_LICENSE from .license.env
SET CLUSTER SETTING enterprise.license = '<key>';
```

(No `cluster.organization` — training keys are bound to an empty organization name, and setting one invalidates them.)

Everything Parts D and E teach — the cross-cluster drill, the verification, the measured RTO, the runbook — works on the free path.

Create a dataset with something worth losing:

```
scripts/crdb sql <<'SQL'
CREATE DATABASE bank;
USE bank;

CREATE TABLE accounts (
  id      UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  name    STRING NOT NULL,
  balance DECIMAL(12,2) NOT NULL,
  region  STRING NOT NULL
);

CREATE TABLE transfers (
  id      UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  from_id UUID NOT NULL REFERENCES accounts(id),
  to_id   UUID NOT NULL REFERENCES accounts(id),
  amount  DECIMAL(12,2) NOT NULL,
  ts      TIMESTAMPTZ NOT NULL DEFAULT now()
);

INSERT INTO accounts (name, balance, region)
SELECT 'user-' || g, 1000.00, (ARRAY['us-east', 'us-west', 'eu-west'])[1 + g % 3]
FROM generate_series(1, 5000) g;
SQL
```

Then open the SQL shell **in the bank database** — the SQL blocks below assume it:

```
scripts/crdb sql -d bank
```

Tasks

Part A: Full, Incremental & Revision-History Backups (15 min)

1. **Full backup** (free). Production backups are normally taken `AS OF SYSTEM TIME '-10s'` — a slightly stale snapshot that never contends with live traffic — and you will use that form later in the lab. The first one is taken *now*:

```
BACKUP DATABASE bank INTO 'nodelocal://1/backups/bank';
```

⚠ **Why not AS OF SYSTEM TIME here too?** You created `bank` moments ago. A backup as of ten seconds ago would be of a cluster in which the database did not exist yet, and fails with `database "bank" does not exist, or invalid RESTORE timestamp`. Harmless in the lab; worth knowing when you script a backup right after a migration creates a table.

- ❑ **Enterprise:** add revision history so any timestamp inside the window is restorable (this one *is* a historical read, so give the database a moment to be ten seconds old):

```
SELECT pg_sleep(10);
BACKUP DATABASE bank INTO 'nodelocal://1/backups/bank'
  AS OF SYSTEM TIME '-10s'
  WITH revision_history;
```

Without a licence this returns `use of BACKUP with revision_history requires an enterprise license` — run the free form above and continue.

2. Inspect without restoring:

```
SHOW BACKUPS IN 'nodelocal://1/backups/bank';
SHOW BACKUP FROM LATEST IN 'nodelocal://1/backups/bank';
SHOW BACKUP FROM LATEST IN 'nodelocal://1/backups/bank' WITH privileges;
```

Note the `start_time / end_time` columns — that window is your RPO for this backup.

3. Change data, then take an incremental — ☐ Enterprise (`INTO LATEST IN` appends to the most recent full backup in the collection):

```
INSERT INTO transfers (from_id, to_id, amount)
SELECT a.id, b.id, 10.00
FROM (SELECT id FROM accounts LIMIT 500) a,
     (SELECT id FROM accounts LIMIT 1) b;

UPDATE accounts SET balance = balance - 10 WHERE region = 'us-east';

BACKUP DATABASE bank INTO LATEST IN 'nodelocal://1/backups/bank' WITH revision_history;
```

No `AS OF SYSTEM TIME` on this one on purpose: you changed the data seconds ago, and a backup as of ten seconds back would capture none of it — the incremental would show 0 rows and look broken. In production, where the backup runs long after the writes, use it. Without a licence, take a second *full* backup into the same collection instead (`BACKUP DATABASE bank INTO 'nodelocal://1/backups/bank';`) and read "one full plus one incremental" below as "two fulls".

4. Read the chain:

```
SHOW BACKUP FROM LATEST IN 'nodelocal://1/backups/bank';
```

One full plus one incremental. The incremental only stores what changed — check the `rows` and `size_bytes` columns to see how much smaller it is.

5. Record a timestamp you can restore back to:

```
SELECT cluster_logical_timestamp() AS marker;
```

Save this value. Then do something regrettable:

```
-- `transfers` has a foreign key onto `accounts`, so the dependent rows go first.
-- (Try it without this line to see the FK error a real "oops" would hit.)
DELETE FROM transfers
WHERE from_id IN (SELECT id FROM accounts WHERE region = 'us-east')
   OR to_id IN (SELECT id FROM accounts WHERE region = 'us-east');

DELETE FROM accounts WHERE region = 'us-east';
SELECT count(*) FROM accounts;
```

The FK is doing its job here. It is also why a partial restore of one table is rarely enough: the tables that reference it have to come back consistently, which is the argument for database- or cluster-level restores over table-level ones.

6. **Take the backup that runs *after* the accident** — **Enterprise**. A point-in-time restore can only target an instant that some revision-history backup *covers*; the marker you saved is later than every backup taken so far, so it is not restorable yet. The next scheduled incremental is what makes it so — and in the lab you are the scheduler:

```
BACKUP DATABASE bank INTO LATEST IN 'nodelocal://1/backups/bank' WITH revision_history;
SHOW BACKUP FROM LATEST IN 'nodelocal://1/backups/bank';
```

The chain now has two incrementals; the second records the deletes as history rather than forgetting the rows. (Free path: nothing to do — Part B step 3 is enterprise-only anyway.)

Part B: Restore — Database, Table, and Point in Time (15 min)

Before you start Part B: a restore gives you back exactly what the backup chain holds. **FROM LATEST** means the state at the *newest* backup — after Part A step 6 that is the post-delete state (3,334 accounts), and only the point-in-time restore in step 3 brings the deleted rows back. On the free path (no step 6) the newest backup predates the delete, so **FROM LATEST** returns all 5,000 — the free-tier version of the same lesson: you can only restore to instants your backups cover. That gap is your RPO; Part C puts a number on it.

1. **Restore the whole database side by side** (never straight over the live one during a drill):

```
RESTORE DATABASE bank FROM LATEST IN 'nodelocal://1/backups/bank'
WITH new_db_name = 'bank_restored';

SELECT count(*) FROM bank_restored.public.accounts; -- 3334: the newest backup is post-
delete
```

2. **Restore a single table:**

```
CREATE DATABASE bank_table_only;
RESTORE TABLE bank.public.accounts
  FROM LATEST IN 'nodelocal://1/backups/bank'
  WITH into_db = 'bank_table_only';

SELECT count(*) FROM bank_table_only.public.accounts;
```

Restoring `transfers` alone would fail — it has a foreign key into `accounts`. Either restore both tables together or add `WITH skip_missing_foreign_keys`.

3. **Point-in-time restore using the marker from Part A** — `Enterprise` (it reads the revision history, and it needs the post-delete incremental from Part A step 6; without that backup the error is `invalid RESTORE timestamp: supplied backups do not cover requested time`). On the free path, skip to step 4; the deleted rows are recoverable from the full backup you took *before* the delete, which is the free-tier version of the same lesson: your RPO is your backup interval. Paste your marker in place of `<marker value>`:

```
RESTORE DATABASE bank
  FROM LATEST IN 'nodelocal://1/backups/bank'
  AS OF SYSTEM TIME '<marker value>'
  WITH new_db_name = 'bank_pit';

SELECT count(*) FROM bank_pit.public.accounts WHERE region = 'us-east'; -- 1666: back
```

The deleted rows are back. `revision_history` is what makes any timestamp within the backup window restorable, not just the backup instants.

4. **Time the restore — this is your measured RTO:**

```
SELECT job_id, job_type, status, finished - started AS duration
FROM [SHOW JOBS]
WHERE job_type = 'RESTORE' ORDER BY created DESC LIMIT 5;
```

Dataset size	Restore duration	Implied restore rate
--------------	------------------	----------------------

Extrapolate: at that rate, how long does your 2 TB production database take? That number is your **actual** RTO, and it is usually much larger than the one in the runbook.

Part D: Cross-Cluster DR Drill (25 min)

The drill that matters: cluster A is gone, bring the data up on cluster B.

1. **Start the standby cluster.** It is the same compose stack under a second project name, on different ports (SQL 26357, console 8180), so both clusters run side by side:

```
CRDB_COMPOSE=docker/labs-b.yml scripts/crdb up
```

```
scripts/crdb ps                                # cluster A
CRDB_COMPOSE=docker/labs-b.yml scripts/crdb ps  # cluster B
```

Both clusters mount the same backup volume at `/backups` (node 1 of each, via `--external-io-dir`). That is deliberate: it stands in for the cloud bucket both clusters would share in production, and it means the drill exercises **restore** rather than file copying.

2. **Take a full cluster backup on A** — cluster backups include users, roles, and settings, which a database backup does not:

```
scripts/crdb sql -e "
CREATE USER app_user;
GRANT CONNECT ON DATABASE bank TO app_user;"
```

```
# `AS OF SYSTEM TIME '-10s'` reads the cluster as it was ten seconds ago – which is
# before you created app_user. Wait the window out, or the restore comes back without
# the user you just added and step 5 looks like a restore bug.
sleep 12
scripts/crdb sql -e "BACKUP INTO 'nodelocal://1/dr/cluster' AS OF SYSTEM TIME '-10s';"
```

A historical backup is historical about *everything*, users and grants included. This is the same trap in production: back up "as of ten minutes ago" right after an emergency grant, and the grant is not in the backup. ☐With a licence, add `WITH revision_history`.

3. **Confirm cluster B can see the backup** — no copying required:

```
CRDB_COMPOSE=docker/labs-b.yml scripts/crdb sql -e \
"SHOW BACKUPS IN 'nodelocal://1/dr/cluster';"
```

4. **Restore onto B and time it:**

```
time CRDB_COMPOSE=docker/labs-b.yml scripts/crdb sql -e \
"RESTORE FROM LATEST IN 'nodelocal://1/dr/cluster';"
```

A full-cluster `RESTORE` must run into a cluster with no user data — which is exactly the DR situation. If B already has the `bank` database from an earlier attempt, reset it with `CRDB_COMPOSE=docker/labs-b.yml scripts/crdb reset`.

5. **Verify the restore, don't assume it:**

```
CRDB_COMPOSE=docker/labs-b.yml scripts/crdb sql -e "  
SHOW DATABASES;  
SELECT count(*) AS accounts FROM bank.public.accounts;  
SELECT count(*) AS transfers FROM bank.public.transfers;  
SELECT sum(balance) AS total_balance FROM bank.public.accounts;"
```

```
CRDB_COMPOSE=docker/labs-b.yml scripts/crdb sql -e "SHOW USERS;"
```

6. **Compare against A** — a DR drill without a comparison is theatre:

```
echo "A:"; scripts/crdb sql --format=tsv -e \  
"SELECT count(*), sum(balance) FROM bank.public.accounts;"  
echo "B:"; CRDB_COMPOSE=docker/labs-b.yml scripts/crdb sql --format=tsv -e \  
"SELECT count(*), sum(balance) FROM bank.public.accounts;"
```

The row count **and** the business checksum must match. One without the other proves nothing.

7. **Fill in the drill record:**

Metric	Value
Backup size	
Backup duration	
Copy/transfer duration	
Restore duration	
Total measured RTO	
Row-count match A vs B	☐ / ☐
Checksum (sum of balances) match	☐ / ☐
Users and grants restored	☐ / ☐

Part E: The Runbook (10 min)

Write the runbook for this drill. It should be short enough to follow while adrenaline is high.

DR Runbook – Restore `bank` to a Standby Cluster

****Trigger:**** primary cluster unavailable for > 10 minutes, or data corruption confirmed.
****Owner:**** on-call DBA. ****Escalation:**** platform lead after 30 minutes.

0. Decide (2 min)

- [] Is this a full-cluster loss (→ full restore) or a bad write (→ point-in-time restore)?
- [] Note the last known-good timestamp. ``SELECT cluster_logical_timestamp();`` if reachable.

1. Locate the backup (2 min)

```
SHOW BACKUPS IN 's3://prod-backups/bank';  
SHOW BACKUP FROM LATEST IN 's3://prod-backups/bank';
```

- [] Confirm ``end_time`` is within RPO.

2. Restore (RT0: measured ____ min per TB)

```
RESTORE FROM LATEST IN 's3://prod-backups/bank' AS OF SYSTEM TIME '<ts>';
```

- [] Watch: `SELECT status, fraction_completed FROM [SHOW JOBS] WHERE job_type='RESTORE';`

3. Verify BEFORE cutting traffic over

- [] Row counts match the last known-good figures
- [] Business checksum matches (e.g. `SUM(balance)`)
- [] Users, roles, and grants present
- [] Application smoke test passes against the standby

4. Cut over

- [] Update connection string / DNS
- [] Confirm error rate returns to baseline
- [] Announce recovery, note actual RT0/RPO achieved

5. Post-incident

- [] Re-establish backup schedules on the new primary
- [] File the gap between measured and documented RT0

Swap runbooks with another pair. Can they follow yours without asking you a question?

Optional — If Time Allows

These parts are not required to complete the lab; they extend it by about 10 minutes. Do them if you finish early, or after class — the cluster and data from the core parts carry over.

Part C: Scheduled Backups & RPO Math (10 min)

1. Create a schedule:

```
-- With a licence: incrementals every 5 minutes, a full every day  
CREATE SCHEDULE bank_backups  
  FOR BACKUP DATABASE bank INTO 'nodelocal://1/backups/scheduled'  
  RECURRING '*/5 * * * *'  
  FULL BACKUP '@daily'  
  WITH SCHEDULE OPTIONS first_run = 'now';
```

On the free path this succeeds but prints:

Without an enterprise license, this schedule will only run full backups. Make that explicit rather than leaving a surprise in the notice:

```
CREATE SCHEDULE bank_backups_full
FOR BACKUP DATABASE bank INTO 'nodelocal://1/backups/scheduled'
RECURRING '@hourly'
FULL BACKUP ALWAYS
WITH SCHEDULE OPTIONS first_run = 'now';
```

This is a real capacity-planning point, not a licensing footnote. Full-only backups make every cycle cost a complete copy — so on the free tier your practical RPO is bounded by how often you can afford a full backup, not by how often you'd like one.

2. Inspect it:

```
SHOW SCHEDULES;
SELECT id, label, schedule_status, next_run, recurrence FROM [SHOW SCHEDULES]
WHERE label LIKE 'bank%';
```

Two rows appear: one for the full-backup schedule and one for the incremental.

3. Watch it fire, then check the jobs it created:

```
SELECT job_id, status, created, description
FROM [SHOW JOBS] WHERE job_type = 'BACKUP' ORDER BY created DESC LIMIT 5;
```

4. The RPO math. With incrementals every 5 minutes:

worst-case RPO = incremental interval + backup duration + detection time

If the incremental takes 40 s and it takes you 2 minutes to notice the outage, your true RPO is roughly $5 \text{ min} + 40 \text{ s} + 2 \text{ min} \approx 7.7 \text{ min}$ — not 5 minutes.

Cadence	Nominal RPO	Realistic RPO	Storage cost
Hourly incremental	1 h		low
15-minute incremental	15 m		medium
5-minute incremental	5 m		high
Continuous (changefeed to object store)	seconds		highest

5. Pause and drop the schedule so it stops firing for the rest of the lab:

```
PAUSE SCHEDULES SELECT id FROM [SHOW SCHEDULES] WHERE label LIKE 'bank%';
DROP SCHEDULES SELECT id FROM [SHOW SCHEDULES] WHERE label LIKE 'bank%';
```

Cleanup

```
scripts/crdb down
CRDB_COMPOSE=docker/labs-b.yml scripts/crdb down
```

That also removes the shared `/backups` volume, so the next run of this lab starts clean.

Lab 11 Deliverables

□ **Backups**: full (and, with a licence, incremental + revision history), inspected via `SHOW BACKUP` □ **Restores**: whole database, single table, and point-in-time into new names □ **Schedule**: created, observed firing, and dropped; RPO computed honestly □ **DR drill**: cross-cluster restore completed and *verified* against the source □ **Measured RTO**: a real number, extrapolated to production data size □ **Runbook**: written, and readable by someone else at 3am

Challenge Exercises

1. **Locality-aware backups**. Take a backup with `BACKUP INTO ('s3://us-east/...?COCKROACH_LOCALITY=region%3Dus-east1', 's3://eu/...?COCKROACH_LOCALITY=default')`. What problem does this solve, and which regulation usually motivates it?
2. **Restore a table that has dropped since backup time**. Use `SHOW BACKUP` to find its name at that revision, then restore only it. What breaks if the schema has changed since?
3. **Backup a subset**. `BACKUP TABLE bank.accounts, bank.transfers INTO ...` and compare size and duration against the database backup. When is table-level worth the complexity?
4. **Measure the impact of revision_history**. Take the same backup with and without it. Compare size and duration. What did the extra bytes buy you?

Reference

Command	Purpose
<code>BACKUP INTO 'sink' AS OF SYSTEM TIME '-10s'</code>	Full backup from a stale snapshot
<code>BACKUP ... INTO LATEST IN 'sink'</code>	Incremental against the latest full
<code>WITH revision_history</code>	Enables arbitrary point-in-time restore in the window
<code>SHOW BACKUPS IN 'sink'</code>	List backup chains at a destination
<code>SHOW BACKUP FROM LATEST IN 'sink'</code>	Contents and timestamps without restoring
<code>RESTORE ... WITH new_db_name = 'x'</code>	Side-by-side restore, non-destructive

Command	Purpose
<code>RESTORE TABLE ... WITH into_db = 'x'</code>	Single-table restore
<code>RESTORE FROM LATEST IN 'sink'</code>	Full-cluster restore (users, roles, settings)
<code>CREATE SCHEDULE ... RECURRING ... FULL BACKUP ...</code>	Managed backup cadence
<code>SELECT cluster_logical_timestamp()</code>	Capture a restorable marker